

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
sns.set_theme(color_codes=True)
pd.set_option('display.max_columns', None)
```

```
df = pd.read_csv('supply_chain_data.csv')
df.head()
```

	Product type	SKU	Price	Availability	Number of products sold	Revenue generated	Customer demographics	Stock levels
0	haircare	SKU0	69.808006	55	802	8661.996792	Non-binary	58
1	skincare	SKU1	14.843523	95	736	7460.900065	Female	53
2	haircare	SKU2	11.319683	34	8	9577.749626	Unknown	1
3	skincare	SKU3	61.163343	68	83	7766.836426	Non-binary	23
4	skincare	SKU4	4.805496	26	871	2686.505152	Non-binary	5

## ▼ Data Preprocessing Part 1

```
df.dtypes
```

```
Product type      object
SKU               object
Price            float64
Availability       int64
Number of products sold  int64
Revenue generated  float64
Customer demographics  object
Stock levels      int64
Lead times        int64
Order quantities  int64
Shipping times    int64
Shipping carriers  object
Shipping costs    float64
Supplier name     object
Location          object
Lead time         int64
Production volumes  int64
Manufacturing lead time  int64
Manufacturing costs  float64
```

```

Inspection results    object
Defect rates          float64
Transportation modes  object
Routes                object
Costs                 float64
dtype: object

```

```
df.select_dtypes(include='object').nunique()
```

```

Product type          3
SKU                   100
Customer demographics 4
Shipping carriers      3
Supplier name         5
Location              5
Inspection results    3
Transportation modes  4
Routes                3
dtype: int64

```

```
df.shape
```

```
(100, 24)
```

```
#Drop SKU Column because this is just supply chain id
```

```
df.drop(columns=['SKU'], inplace=True)
```

```
df.shape
```

```
(100, 23)
```

```
df.nunique()
```

```

Product type          3
Price                 100
Availability           63
Number of products sold 96
Revenue generated     100
Customer demographics  4
Stock levels          65
Lead times            29
Order quantities       61
Shipping times         10
Shipping carriers      3
Shipping costs         100
Supplier name          5
Location               5
Lead time              29
Production volumes     96
Manufacturing lead time 30
Manufacturing costs    100
Inspection results     3
Defect rates           100
Transportation modes   4
Routes                 3
Costs                  100
dtype: int64

```

## ▼ Exploratory Data Analysis


```
# list of categorical variables to plot
cat_vars = ['Product type', 'Customer demographics', 'Shipping carriers', 'Supplier name', 'Location',
            'Inspection results', 'Transportation modes', 'Routes']

# create figure with subplots
fig, axs = plt.subplots(nrows=2, ncols=4, figsize=(20, 10))
axs = axs.flatten()

# create barplot for each categorical variable
for i, var in enumerate(cat_vars):
    sns.barplot(x=var, y='Costs', data=df, ax=axs[i], estimator=np.mean)
    axs[i].set_xticklabels(axs[i].get_xticklabels(), rotation=90)

# adjust spacing between subplots
fig.tight_layout()

# show plot
plt.show()
```



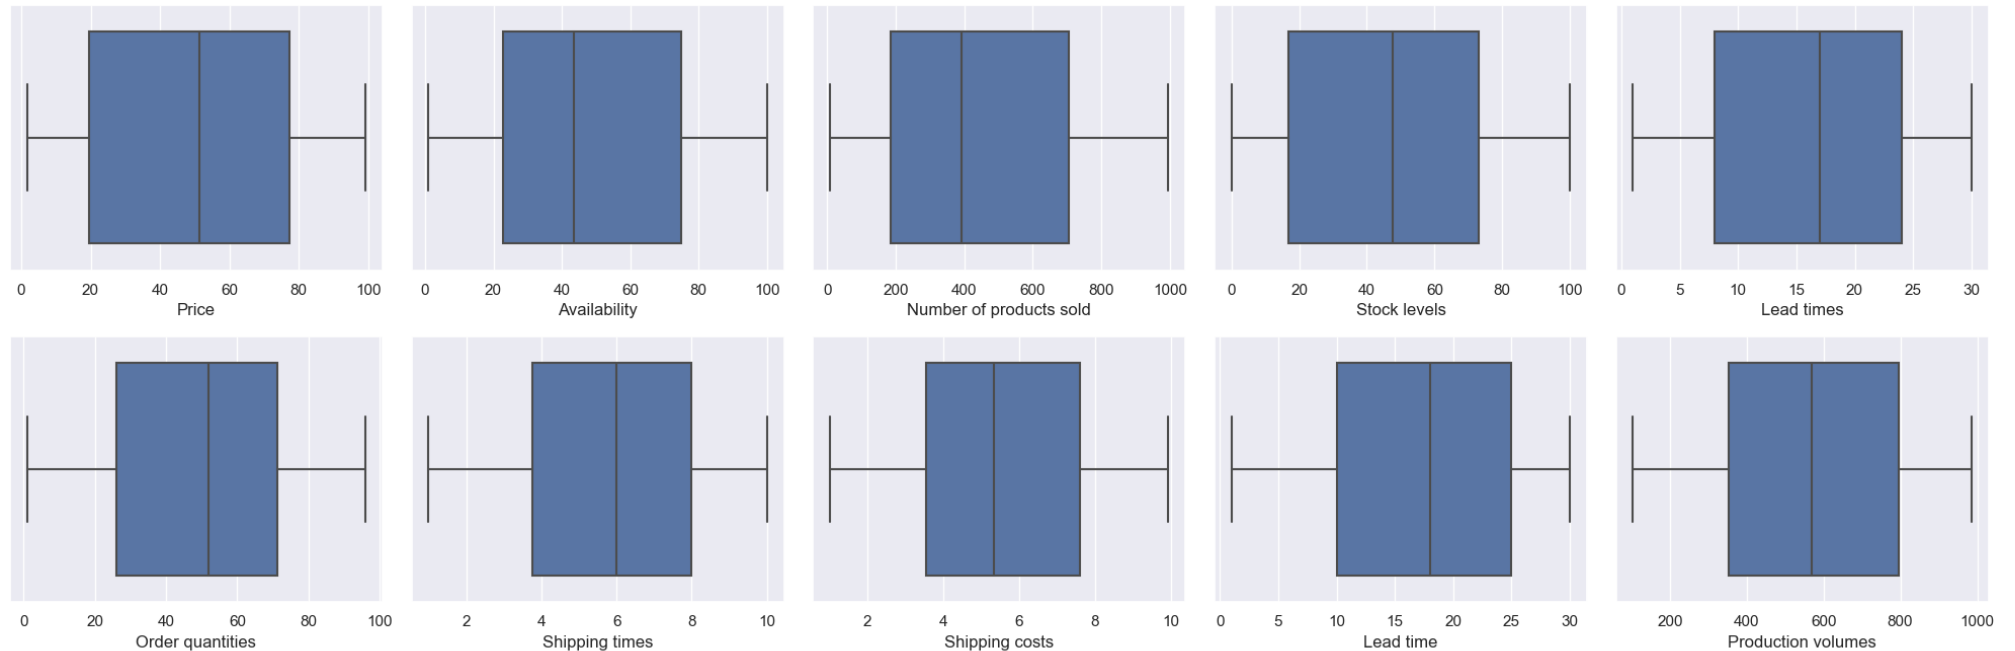
```
num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
            'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
            'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
            'Defect rates']
```

```
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()
```

```
for i, var in enumerate(num_vars):
    sns.boxplot(x=var, data=df, ax=axs[i])
```

```
# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])
```

```
fig.tight_layout()
plt.show()
```



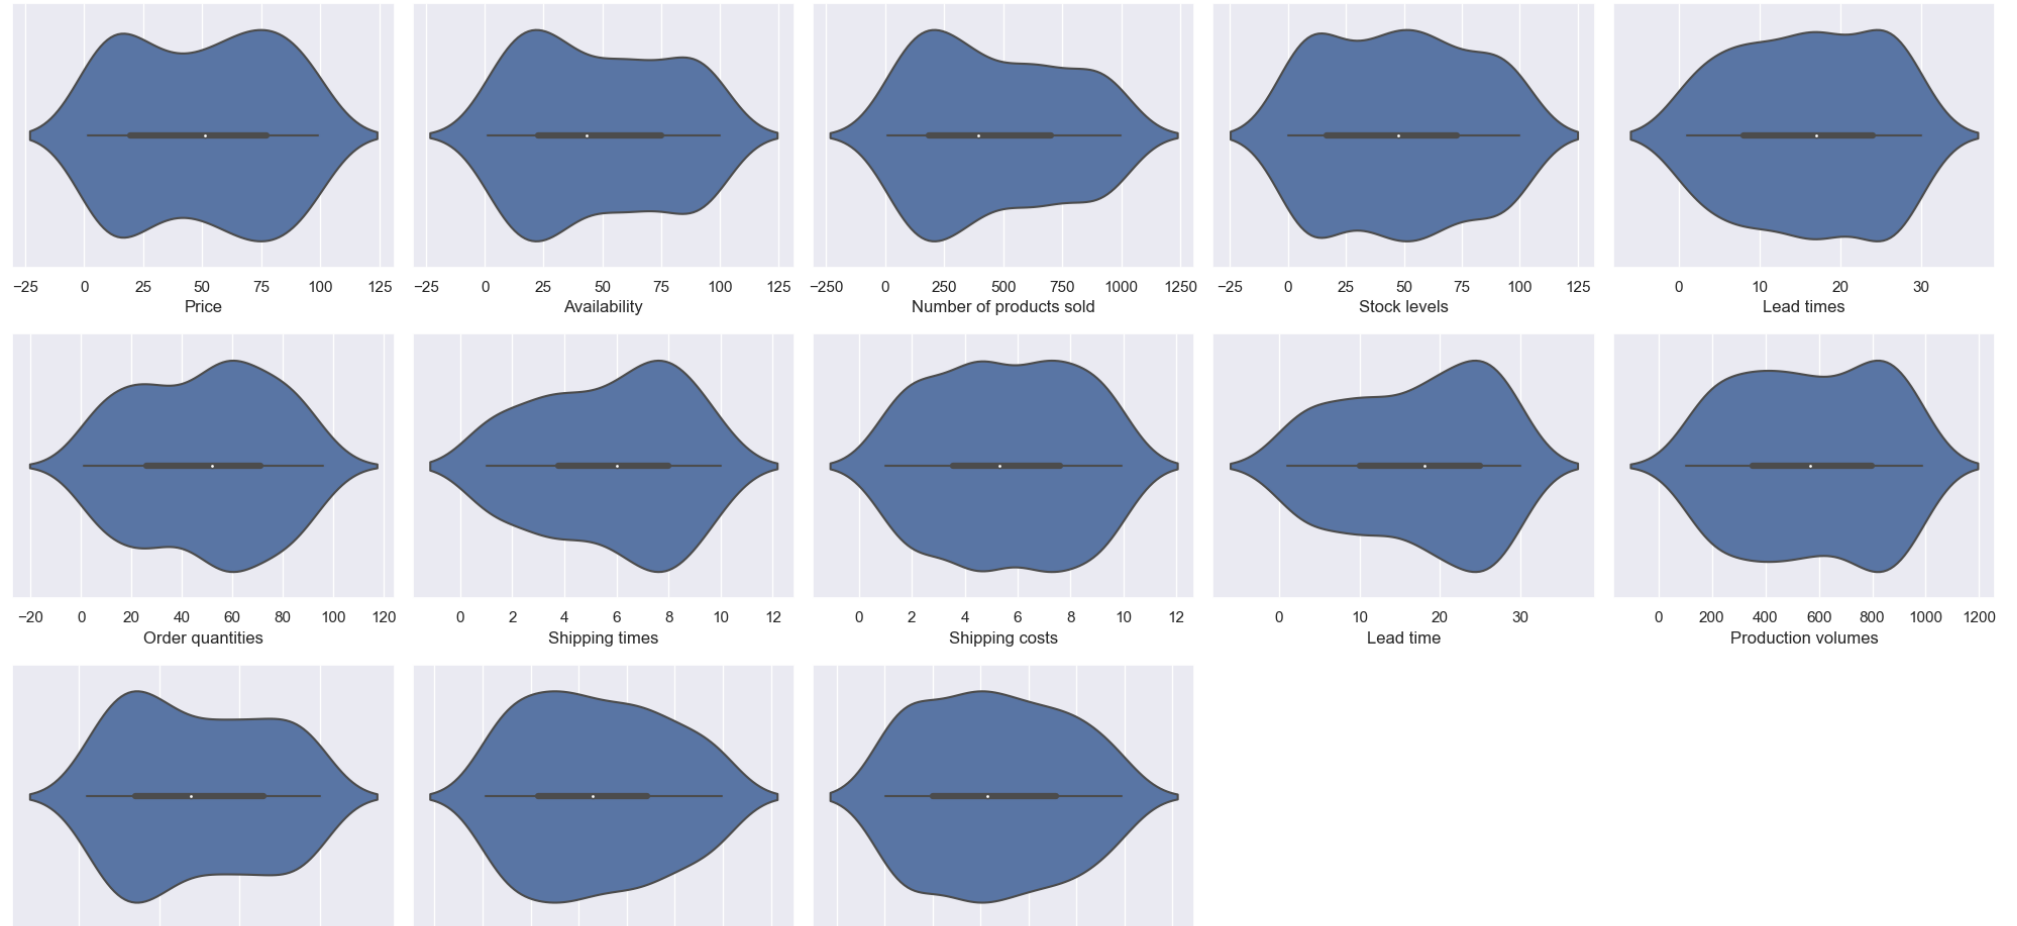
```
num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
            'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
            'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
            'Defect rates']
```

```
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()
```

```
for i, var in enumerate(num_vars):
    sns.violinplot(x=var, data=df, ax=axs[i])
```

```
# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])
```

```
fig.tight_layout()
plt.show()
```



```
num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
            'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
            'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
            'Defect rates']
```

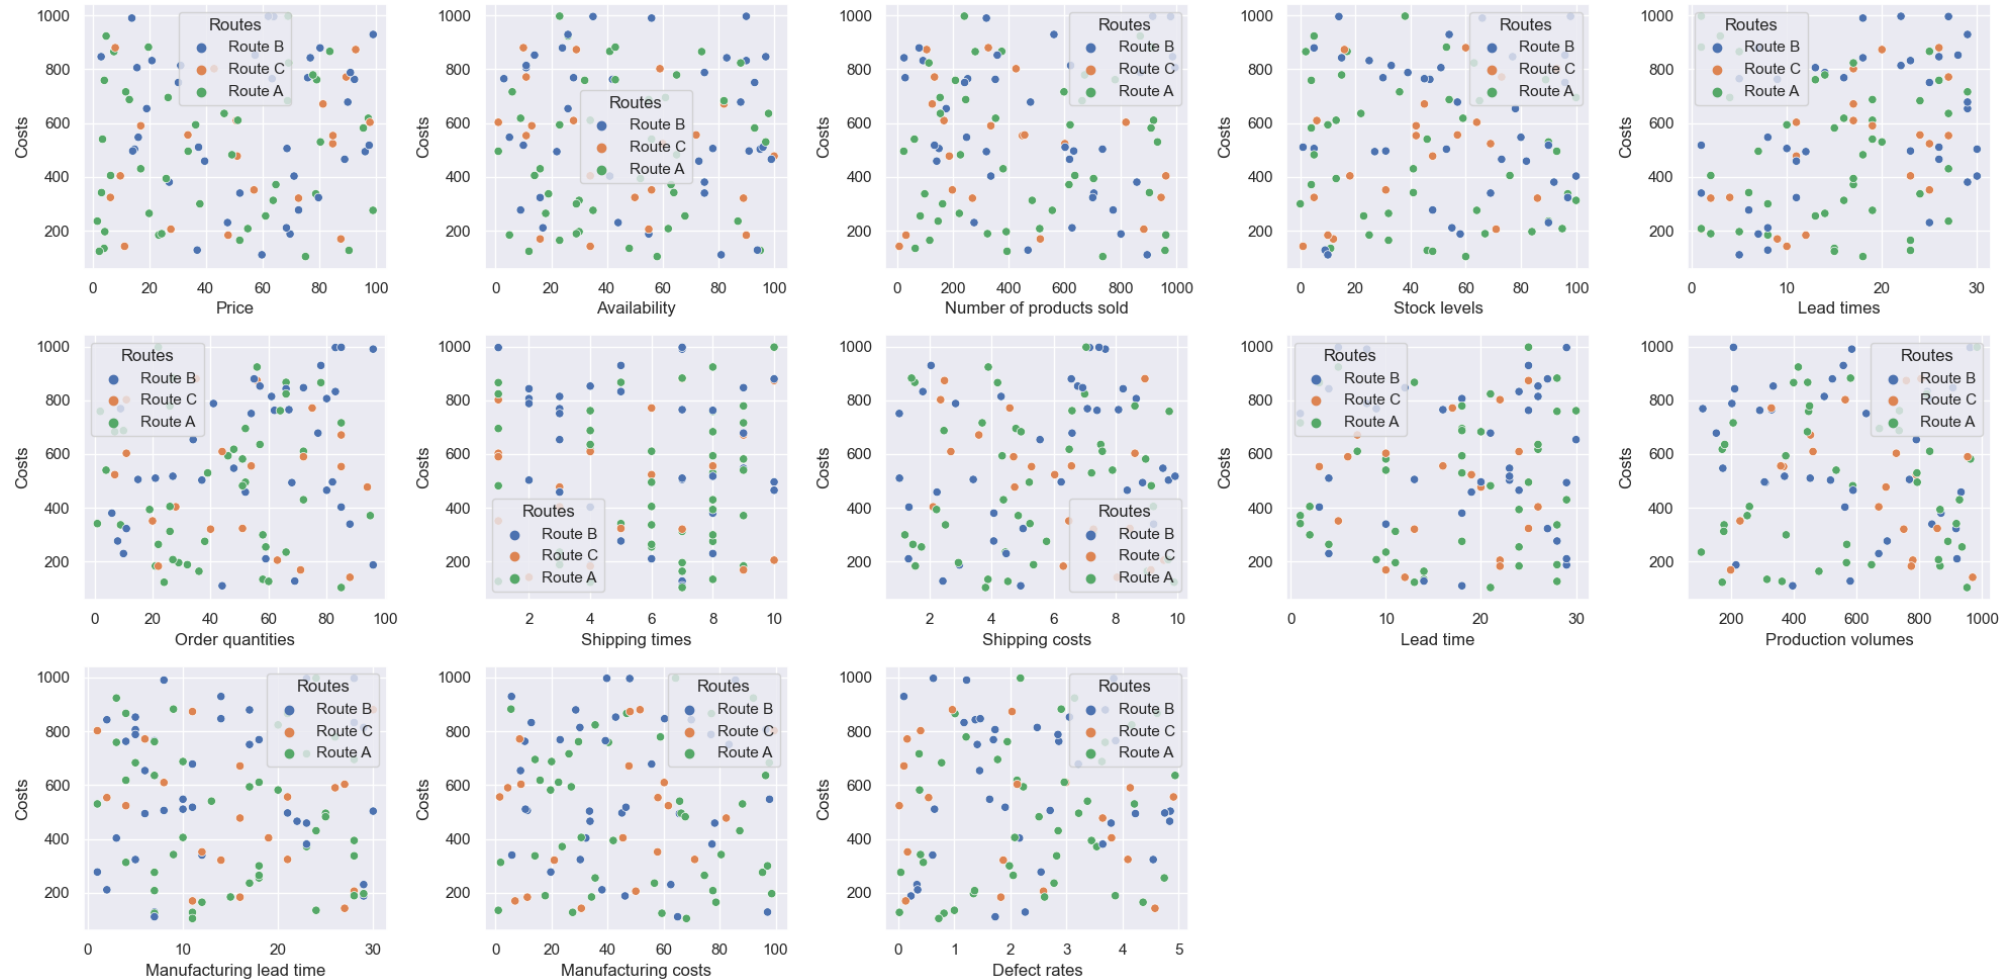
```
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()
```

```
for i, var in enumerate(num_vars):
    sns.scatterplot(x=var, y='Costs', hue='Routes', data=df, ax=axs[i])
```

```
# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])
```

```
fig.tight_layout()
```

```
plt.show()
```



```
num_vars = ['Price', 'Availability', 'Number of products sold', 'Stock levels',
            'Lead times', 'Order quantities', 'Shipping times', 'Shipping costs',
            'Lead time', 'Production volumes', 'Manufacturing lead time', 'Manufacturing costs',
            'Defect rates']
```


```
fig, axs = plt.subplots(nrows=3, ncols=5, figsize=(20, 10))
axs = axs.flatten()

for i, var in enumerate(num_vars):
    sns.histplot(x=var, data=df, ax=axs[i])

# remove the 14th subplot
fig.delaxes(axs[13])
# remove the 15th subplot
fig.delaxes(axs[14])

fig.tight_layout()
plt.show()
```





## ▼ Data Preprocessing Part 2

```
#Check the missing value
check_missing = df.isnull().sum() * 100 / df.shape[0]
check_missing[check_missing > 0].sort_values(ascending=False)
```

```
Series([], dtype: float64)
```



## ▼ Label Encoding for Object datatypes

```
# Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:
```

```
# Print the column name and the unique values
print(f"{col}: {df[col].unique()}")
```

```
Product type: ['haircare' 'skincare' 'cosmetics']
Customer demographics: ['Non-binary' 'Female' 'Unknown' 'Male']
Shipping carriers: ['Carrier B' 'Carrier A' 'Carrier C']
Supplier name: ['Supplier 3' 'Supplier 1' 'Supplier 5' 'Supplier 4' 'Supplier 2']
Location: ['Mumbai' 'Kolkata' 'Delhi' 'Bangalore' 'Chennai']
Inspection results: ['Pending' 'Fail' 'Pass']
Transportation modes: ['Road' 'Air' 'Rail' 'Sea']
Routes: ['Route B' 'Route C' 'Route A']
```

```
from sklearn import preprocessing
```

```
# Loop over each column in the DataFrame where dtype is 'object'
for col in df.select_dtypes(include=['object']).columns:
```

```
# Initialize a LabelEncoder object
label_encoder = preprocessing.LabelEncoder()
```

```
# Fit the encoder to the unique values in the column
label_encoder.fit(df[col].unique())
```

```
# Transform the column using the encoder
df[col] = label_encoder.transform(df[col])
```

```
# Print the column name and the unique encoded values
print(f"{col}: {df[col].unique()}")
```

```
Product type: [1 2 0]
Customer demographics: [2 0 3 1]
Shipping carriers: [1 0 2]
```

```

Supplier name: [2 0 4 3 1]
Location: [4 3 2 0 1]
Inspection results: [2 0 1]
Transportation modes: [2 0 1 3]
Routes: [1 2 0]

```

df.dtypes

```

Product type          int32
Price                 float64
Availability           int64
Number of products sold int64
Revenue generated      float64
Customer demographics int32
Stock levels          int64
Lead times            int64
Order quantities       int64
Shipping times         int64
Shipping carriers      int32
Shipping costs         float64
Supplier name          int32
Location              int32
Lead time             int64
Production volumes     int64
Manufacturing lead time int64
Manufacturing costs    float64
Inspection results     int32
Defect rates           float64
Transportation modes   int32
Routes                int32
Costs                 float64
dtype: object

```

There's no outlier so we dont have to remove it

## ▼ Correlation Heatmap

```

#Correlation Heatmap
plt.figure(figsize=(20, 16))
sns.heatmap(df.corr(), fmt='.2g', annot=True)

```

&lt;AxesSubplot:&gt;



## ▼ Train test Split

```
X = df.drop('Costs', axis=1)
y = df['Costs']

#test size 20% and train size 80%
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,random_state=0)
```

## ▼ Decision Tree Regressor

```
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import load_boston

# Create a DecisionTreeRegressor object
dtree = DecisionTreeRegressor()

# Define the hyperparameters to tune and their values
param_grid = {
    'max_depth': [2, 4, 6, 8],
    'min_samples_split': [2, 4, 6, 8],
    'min_samples_leaf': [1, 2, 3, 4],
    'max_features': ['auto', 'sqrt', 'log2'],
    'random_state': [0, 7, 42]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(dtree, param_grid, cv=5, scoring='neg_mean_squared_error')

# Fit the GridSearchCV object to the data
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print(grid_search.best_params_)

{'max_depth': 2, 'max_features': 'sqrt', 'min_samples_leaf': 3, 'min_samples_split': 2, 'random_state': 0}

from sklearn.tree import DecisionTreeRegressor
dtree = DecisionTreeRegressor(random_state=0, max_depth=2, max_features='sqrt', min_samples_leaf=3, min_samples_split=2)
dtree.fit(X_train, y_train)

DecisionTreeRegressor(max_depth=2, max_features='sqrt', min_samples_leaf=3,
                      random_state=0)
```

```
from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = dtree.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)


print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))

MAE is 248.4413893861546
MAPE is 0.5893818876444419
MSE is 72806.47766651674
R2 score is -0.08647889188367719
RMSE score is 269.8267549123266


imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": dtree.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (Decision Tree Regressor)', fontsize=18)
plt.xlabel ('Importance', fontsize=16)
plt.ylabel ('Feature Name', fontsize=16)
plt.show()
```

## Feature Importance Each Attributes (Decision Tree Regressor)



```
import shap
explainer = shap.TreeExplainer(dtree)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```



## ▼ Random Forest Regressor

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV

# Create a Random Forest Regressor object
rf = RandomForestRegressor()

# Define the hyperparameter grid
param_grid = {
    'max_depth': [3, 5, 7, 9],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt'],
    'random_state': [0, 7, 42]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='r2')

# Fit the GridSearchCV object to the training data
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print("Best hyperparameters: ", grid_search.best_params_)

# SHAP value (impact on model output)
from sklearn.ensemble import RandomForestRegressor
rf = RandomForestRegressor(random_state=0, max_depth=3, min_samples_split=5, min_samples_leaf=2,
                           max_features='sqrt')
rf.fit(X_train, y_train)

RandomForestRegressor(max_depth=3, max_features='sqrt', min_samples_leaf=2,
                       min_samples_split=5, random_state=0)
```

```

from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = rf.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

```

```

print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))

```

```

MAE is 247.33969719962744
MAPE is 0.6029768224226728
MSE is 71899.28833186119
R2 score is -0.07294105713825938
RMSE score is 268.14042651540103

```

```

imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": dtree.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

```


```

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (Random Forest Regressor)', fontsize=18)
plt.xlabel ('Importance', fontsize=16)
plt.ylabel ('Feature Name', fontsize=16)
plt.show()


```



## Feature Importance Each Attributes (Random Forest Regressor)



```
import shap
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(X_test)
shap.summary_plot(shap_values, X_test)
```



## AdaBoost Regressor

```

Production volumes
from sklearn.ensemble import AdaBoostRegressor
from sklearn.model_selection import GridSearchCV

# Create an AdaBoost Regressor object
ada = AdaBoostRegressor()

# Define the hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 150, 200],
    'learning_rate': [0.01, 0.1, 1],
    'loss': ['linear', 'square', 'exponential'],
    'random_state': [0, 7, 42]
}

# Create a GridSearchCV object
grid_search = GridSearchCV(ada, param_grid, cv=5, scoring='r2')

# Fit the GridSearchCV object to the training data
grid_search.fit(X_train, y_train)

# Print the best hyperparameters
print("Best hyperparameters: ", grid_search.best_params_)

Best hyperparameters: {'learning_rate': 1, 'loss': 'linear', 'n_estimators': 50, 'random_state': 7}

# SHAP values (import as model output)
from sklearn.ensemble import AdaBoostRegressor
ada = AdaBoostRegressor(random_state=7, n_estimators=50, learning_rate=1, loss='linear')
ada.fit(X_train, y_train)

AdaBoostRegressor(learning_rate=1, random_state=7)

```


```
from sklearn import metrics
from sklearn.metrics import mean_absolute_percentage_error
import math
y_pred = ada.predict(X_test)
mae = metrics.mean_absolute_error(y_test, y_pred)
mape = mean_absolute_percentage_error(y_test, y_pred)
mse = metrics.mean_squared_error(y_test, y_pred)
r2 = metrics.r2_score(y_test, y_pred)
rmse = math.sqrt(mse)

print('MAE is {}'.format(mae))
print('MAPE is {}'.format(mape))
print('MSE is {}'.format(mse))
print('R2 score is {}'.format(r2))
print('RMSE score is {}'.format(rmse))


MAE is 255.28612180541396
MAPE is 0.5859844238678936
MSE is 78800.74415400976
R2 score is -0.17593032834538103
RMSE score is 280.71470241868303

imp_df = pd.DataFrame({
    "Feature Name": X_train.columns,
    "Importance": ada.feature_importances_
})
fi = imp_df.sort_values(by="Importance", ascending=False)

fi2 = fi.head(10)
plt.figure(figsize=(10,8))
sns.barplot(data=fi2, x='Importance', y='Feature Name')
plt.title('Feature Importance Each Attributes (AdaBoost Regressor)', fontsize=18)
plt.xlabel ('Importance', fontsize=16)
plt.ylabel ('Feature Name', fontsize=16)
plt.show()
```



All of the Algorithms got bad R2 Score and MAPE Score even with hyperparameter tuning because we only have 100 data and the distribution is spread



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.